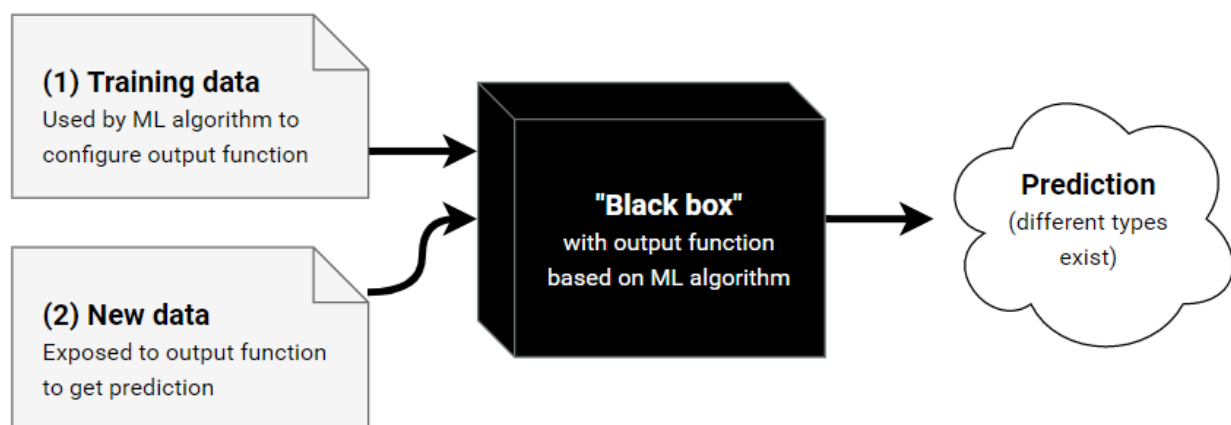


- Lecture 2
 - Linear Regression Models
 - Cat Hearts example:
 - Experience E
 - Learning Task, T
 - Linear Regression Model
 - Performance Measure, P
 - Unconstrained Optimisation (Minimisation)
 - Differentiation Rules
 - Approach 1: Ordinary least squares
 - Approach 2: Gradient descent
 - Attempt 1 (**failed**)
 - Attempt 2: Gradient Descent (1D)

Lecture 2

Linear Regression Models



Cat Hearts example:

Experience E

- The dataset consists of n data points
 - $((x_1, y_1), \dots, (x_n, y_n)) \in \mathbb{R}^d \times \mathbb{R}$
 - $x_i \in \mathbb{R}^d$ is the "input" for the i^{th} data point as a feature vector with d elements, d being the # of dimensions in the feature space, in this case 1.
 - $y_i \in \mathbb{R}$ is the "output" for the i^{th} data point, in this case the weight of the corresponding cat heart.

Learning Task, T

- In this example, our task is: **Linear Regression**

- Find a "model", i.e. a function:
 - $f : \mathbb{R}^d \rightarrow \mathbb{R}$
- s.t. our future observations produce output "close to" the true output.

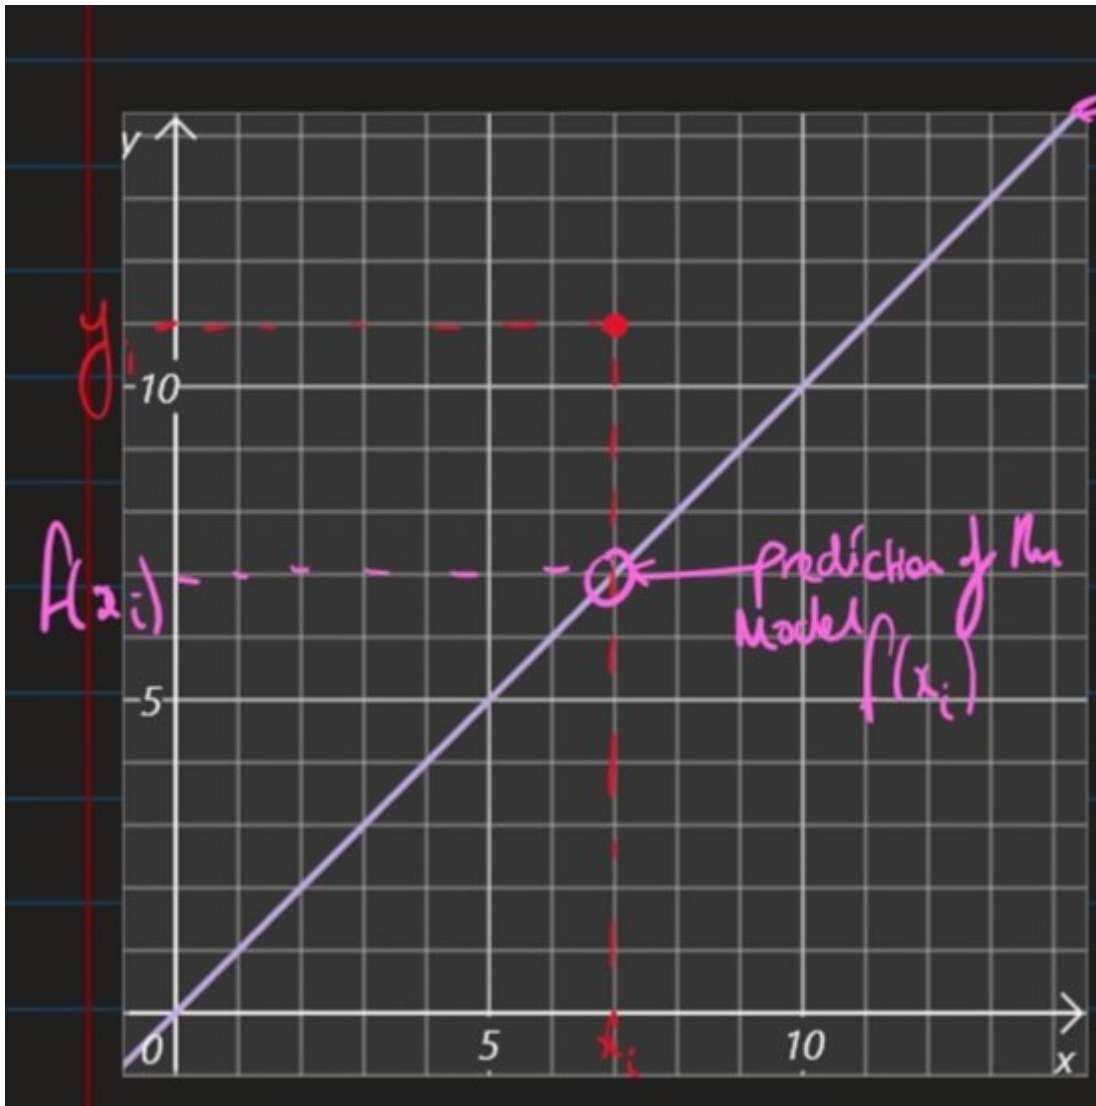
Linear Regression Model

- A linear regression model has the form:
 - $f(x) = (\sum_{i=1}^d w_i \cdot x_i) + b$
 - where:
 - $x \in \mathbb{R}^d$ is the input vector (feature)
 - $w \in \mathbb{R}^d$ is the weight vector (parameters)
 - $b \in \mathbb{R}$ is a bias (parameter)
 - $f(x) \in \mathbb{R}$ is the predicted output
-

- In our cat example we have:
 - $d=1$ as "body weight" is our only feature
 - $b=0$ as from intuition we expect a cat of 0 weight to have a heart of 0 weight.
 - Our model has one parameter: w

Performance Measure, $J(w)$

- Want a function, $J(w)$ which quantifies the error in the predictions for a given parameter w




- The following empirical loss function, J takes into account the errors $\forall n$ data points.
 - $J(w) = (1/2N) \sum_{i=1}^N (y_i - wx_i)^2$
 - where the summation term is squared so that:
 - we ignore the sign
 - we penalise large errors more
- To find the optimum weight, solve:
 - $\frac{\partial J}{\partial w} = 0$

Unconstrained Optimisation (Minimisation)

Given a continuous function:

- $f: \mathbb{R}^d \rightarrow \mathbb{R}$, as our *loss function*
- an element $x \in \mathbb{R}^d$ is called:
 - A **global** minimum of f iff:
 - $\forall y \in \mathbb{R}^d, f(x) \leq f(y)$
 - A **local** minimum of f iff:
 - $\exists \epsilon > 0, \forall y \in \mathbb{R}^d$ if $\forall i \in \{1, \dots, d\}, |x_i - y_i| < \epsilon$ implies $f(x) \leq f(y)$



Theorem: For any continuous function, $f: \mathbb{R} \rightarrow \mathbb{R}$, if x is a local optimum, $f'(x) = 0$

Definition: The 1^{st} derivative of a function $f: \mathbb{R} \rightarrow \mathbb{R}$ is $f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$

Differentiation Rules

1. $(cf(x))' = cf'(x)$
2. $(x^k)' = kx^{k-1}$, if $k \neq 0$
3. $(f(x) + g(x))' = f'(x) + g'(x)$
4. $(f(g(x)))' = f'(g(x))g'(x)$ \leftarrow **chain rule**

Approach 1: Ordinary least squares

- Optimise J by solving $J'(w) = 0$
 - $J(w) = \frac{1}{2N} \sum_{i=1}^N (y_i - wx_i)^2$
 - $J'(w) = \frac{1}{N} \sum_{i=1}^N (wx_i - y_i)x_i$
 - $J'(w) = 0$
 - $\frac{1}{N} \sum_{i=1}^N (wx_i - y_i)x_i = 0$
 - $w \sum_{i=1}^N (x_i)^2 = \sum_{i=1}^N x_i y_i$
 - $w = \frac{\sum_{i=1}^N x_i y_i}{\sum_{i=1}^N x_i^2}$
 - This only has one solution \therefore a global minimum.

Approach 2: Gradient descent

- Often difficult / impossible to solve $J'(w) = 0$ for non-linear models with many parameters

Idea:

- Start with an initial guess
- While $J'(w) \neq 0$:
 - move *slightly* in the *right direction*
- To make this viable we need to define:

- "what is the right direction?"
- "what is slightly?"

Attempt 1 (failed)

$w \leftarrow \text{initial_weight}$ repeat: if $J'(w) < 0$ $w \leftarrow w + \epsilon$ elseif $J'(w) > 0$ $w \leftarrow w - \epsilon$

- where ϵ is the learning rate set manually. **(hyper-parameter)**

Issue with this attempt:

- w may oscillate in the interval $[w_{\text{opt}} - \epsilon, w_{\text{opt}} + \epsilon]$
- w fails to converge

Attempt 2: Gradient Descent (1D)

$w \leftarrow \text{initial_weight}$ repeat: if $J'(w) < 0$ $w \leftarrow w - \epsilon \cdot J'(w)$